

Block-masked identity regularisation closes the qft $\rightarrow$ blocked gap

Setting. DIV2K-8q, QFT(8, 8) basis. `blocked_8` reaches 32.26 dB at $\rho = 0.20$; it is a stable local minimum of QFT(8, 8) (the `qft_warmstart_from_trained_blocked` cell verified this). `qft` (analytic init) reaches 31.29, `qft_identity` (identity init) reaches 31.66. Neither generic-init path crosses to the blocked basin in 1008 steps. Can a structural prior bridge the gap?

Construction

Total loss with regulariser:

$$\mathcal{L}_{\text{total}}(\theta) = \mathcal{L}_{\text{MSE-topK}}(\theta) + \lambda \cdot R_{\text{block}}(\theta)$$

$$R_{\text{block}}(\theta) = \sum_{g \in \mathcal{G}_{\text{outer}}} W \cdot \|T_g - I_g\|_F^2 + \sum_{g \in \mathcal{G}_{\text{inner}}} \|T_g - I_g\|_F^2$$

A gate g is **inner** iff every qubit it touches lies in $\{1, 2, 3\}$ (axis 1) or $\{9, 10, 11\}$ (axis 2). Otherwise **outer**. Per-kind identity element: $I_g = I_2$ for Hadamards; $I_g = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix} = \text{controlled_phase_diag}(\theta)$ for CPs.

gate kind	inner	outer	total
Hadamards	6	10	16
controlled-phase	6	50	56
total	12	60	72

The 12/60 split matches `blocked_8`'s sparse shape bit-exactly: 12 trained QFT(3, 3) gates + 60 identity-pinned gates.

Notation

symbol	meaning
θ	72 gate tensors $\in U(2)^{72}$
g	gate label; $\sum_g$ is over the 72 gates
$\mathcal{G}_{\text{inner}}$	12 inner gates
$\mathcal{G}_{\text{outer}}$	60 outer gates ($\mathcal{G}_{\text{inner}} \cap \mathcal{G}_{\text{outer}} = \emptyset$)
$T_g(\theta)$	current 2×2 unitary at gate g
I_g	identity element for gate g 's kind (H or CP)
$\ M\ _F$	Frobenius norm: $\sqrt{\sum_{ij} M_{ij} ^2}$
λ	reg strength; swept $\{0, 10^{-3}, 10^{-2}, 10^{-1}, 1, 10\}$
W	outer-gate weight; fixed $W = 10$

Key invariant

By construction, at the blocked optimum the outer contribution to R_{block} is zero (every outer gate is bit-exactly at its identity element):

$$R_{\text{block}}(\theta_{\text{blocked} \setminus 8}) = \sum_{g \in \mathcal{G}_{\text{inner}}} \|T_g - I_g\|_F^2 + \varepsilon, \quad \varepsilon \approx 0$$

Empirically $R_{\text{block}}(\theta_{\text{blocked} \setminus 8}) \approx 50.3$ at $W = 10$: inner-sum ≈ 34.5 , outer-sum ≈ 1.6 (dominated by one CP outlier; the other 59 outer gates contribute $< 10^{-2}$ each). The reg term does **not** push the operator away from blocked even at large λ — this is the design property that enables crank-up without collapse.

Implementation

`pdft_benchmarks.identity_reg.BlockMaskedIdentityRegQFTMSELoss`, via the `MSELoss._extra_loss` hook (pdft PR #18). Headline preset: 1008 steps, batch 50, val split 0.15, seed 42, `--no-early-stop`. Start from `qft_identity` init.

Result

Test PSNR (dB) after 1008 steps:

basis / reg	$\rho = 0.05$	$\rho = 0.10$	$\rho = 0.15$	$\rho = 0.20$
qft (analytic init)	25.09	27.57	29.53	31.29
qft_identity (identity init)	25.23	27.81	29.84	31.66
qft_identity + reg, $\lambda = 10^{-3}$	25.23	27.81	29.84	31.66
qft_identity + reg, $\lambda = 10^{-2}$	25.23	27.81	29.84	31.66
qft_identity + reg, $\lambda = 10^{-1}$	25.24	27.81	29.84	31.66
qft_identity + reg, $\lambda = 1$	25.18	28.09	30.30	32.26
qft_identity + reg, $\lambda = 10$	25.18	28.08	30.29	32.25
blocked_8 (warm-start ceiling)	25.18	28.09	30.30	32.26

At $\lambda = 1$, the trained operator matches **blocked_8 bit-exactly to 2 decimals at every ρ** . Trajectory crosses from 31.66 dB (qft_identity basin) to 32.26 dB (blocked basin), closing the full 0.60 dB gap. $\lambda = 10$ within 0.01 dB.




Figure 1: Left: test PSNR at $\rho = 0.20$ vs λ (log x). Sharp transition $31.66 \rightarrow 32.26$ dB between $\lambda = 10^{-1}$ and $\lambda = 1$. Right: per-gate $\|T_g - I_g\|_F$ at end of training, inner (12, orange) vs outer (60, blue). At $\lambda \geq 1$, outer gates collapse to a thin line at zero — **all 60 pinned bit-exactly at identity**, matching **blocked_8**’s structural shape. Inner-gate distributions are invariant across λ : the prior leaves them free.

Interpretation

Outcome (i) from the spec’s three-falsifiable-outcomes frame: the blocked basin was **unfavored, not isolated**. Given the matched structural prior, smooth-Riemannian Adam reaches it from identity init in 1008 steps with no curriculum and no warm-start.

Leading-prior caveat. The regulariser bakes in **blocked_8**’s known structural shape. A positive result is “if you tell the optimiser the answer-shape, it finds the answer” — closer to warm-start with extra steps than to basin discovery from scratch.

The shift in framing: the headline **qft < blocked_8** gap is now attributable to **initialisation + optimiser-coupling**, not to a loss-landscape obstruction. The remaining open question is **how weak a prior suffices**.

L1-to-identity: dropping the block-aligned mask

Drop the inner/outer mask; penalise all gates uniformly with an L1 norm (Huber-smoothed to avoid the kink at $T_g = I_g$):

$$R_{L1}(\theta) = \sum_g \sqrt{\|T_g - I_g\|_F^2 + \varepsilon}, \quad \varepsilon = 10^{-12}$$

The unsmoothed L1's gradient $\frac{T_g - I_g}{\|T_g - I_g\|_F}$ is $\frac{0}{0}$ at qft_identity init; `jax.grad` returns NaN there. The ε -smoothing yields a finite ($= 0$) gradient at the kink while agreeing with true L1 to leading order for $\|T_g - I_g\|_F \gtrsim 10^{-5}$. L1 induces sparsity from training dynamics: gates either move appreciably or stay pinned; no continuous shrinkage.

L1 sweep — DIV2K-8q. $\lambda \in \{0.1, 1, 10\}$, qft_identity init, same headline preset:

basis / reg	$\rho = 0.05$	$\rho = 0.10$	$\rho = 0.15$	$\rho = 0.20$	active gates
qft_identity (no reg)	25.23	27.81	29.84	31.66	72/72
L1 reg, $\lambda = 0.1$	25.23	27.81	29.84	31.66	72/72
L1 reg, $\lambda = 1$	25.23	27.81	29.84	31.66	72/72
L1 reg, $\lambda = 10$	25.18	28.08	30.30	32.26	6/72
blocked_8 (ceiling)	25.18	28.09	30.30	32.26	12/72

L1 at $\lambda = 10$ matches **blocked_8** to ≈ 0.02 dB at every ρ , with **only 6 active gates** (vs **blocked_8's** 12). The 6 are exactly the inner-block Hadamards (indices 0, 1, 2, 8, 9, 10); all 6 inner CPs that **blocked_8** makes non-trivial are pinned at identity by L1. The 6 active gates trained to rotation-by- $\frac{\pi}{4}$ matrices $\begin{pmatrix} c & s \\ -s & c \end{pmatrix}$ with $c, s \approx \frac{1}{\sqrt{2}}$ — **not** Hadamards $\begin{pmatrix} c & s \\ s & -c \end{pmatrix}$. The basin is structurally distinct from **blocked_8** (no CP entanglement) yet PSNR-equivalent.

Comparison against the full block-size grid (DIV2K-8q):

basis	$\rho = 0.05$	$\rho = 0.10$	$\rho = 0.15$	$\rho = 0.20$
blocked_4	19.06	26.93	29.76	32.05
blocked_8	25.18	28.09	30.30	32.26
blocked_16	25.23	27.81	29.84	31.66
blocked_32	24.91	27.30	29.20	30.91
L1 $\lambda = 10$	25.18	28.08	30.30	32.26

L1 hits the **training-rate-optimal** blocked ceiling at every ρ (**blocked_8** wins at $\rho \geq 0.10$; at $\rho = 0.05$ **blocked_16** wins by 0.05 dB and L1 matches **blocked_8** rather than **blocked_16**). Loss trains at $\rho = 0.10$ via $k = \lfloor 2^{m+n} \cdot 0.1 \rfloor$.

L1 sweep — QuickDraw ($m = n = 5$). Same procedure, 30-gate QFT(5, 5) basis:

basis / reg	$\rho = 0.05$	$\rho = 0.10$	$\rho = 0.15$	$\rho = 0.20$	active gates
qft (analytic init)	16.72	19.58	22.05	24.36	30/30
blocked_4 / blocked_8	18.12	22.40	26.18	30.04	varies
L1 reg, $\lambda = 0.1$	17.24	23.33	30.06	40.62	2/30
L1 reg, $\lambda = 1$	12.82	18.35	30.48	54.11	0/30
L1 reg, $\lambda = 10$	12.82	18.34	30.85	61.89	0/30

On QuickDraw, L1 finds a **different** answer: at $\lambda \geq 1$ **every gate is pinned at identity** and the trained operator is literally $T(x) = x$. PSNR at $\rho = 0.20$ reaches 61.89 dB — 30 + dB above any **blocked_b** and 37 + dB above analytic **qft**. QuickDraw images are sparse line drawings; top- k truncation in the pixel domain already captures essentially all the strokes. A transform basis is **worse** than no transform.

Caveat. The 61.89 dB QuickDraw number is partly a property of the data, not just the regulariser: at 32×32 with mostly-zero pixels, top-20%-of-pixels retention is near-perfect by construction. The DIV2K result is the more discriminating signal because DIV2K is not pixel-sparse.




Figure 2: Left, middle: test PSNR at $\rho = 0.20$ vs λ for L1-regularised qft_identity on DIV2K-8q and QuickDraw. On DIV2K, L1 matches **blocked_8** at $\lambda = 10$; on QuickDraw L1 **exceeds** every **blocked_b** by 30+ dB at $\lambda \geq 1$. Right: end-of-training per-gate $\|T_g - I_g\|_F$ at $\lambda = 10$ for both datasets, coloured by **blocked_8**'s inner (12, orange) / outer (60, blue) mask. DIV2K: 6 inner-Hadamard gates are active at distance ≈ 1.08 ; everything else is bit-exactly at identity. QuickDraw: **all** gates are bit-exactly at identity — L1 picked no transform.

What this reveals

L1 is **not** finding **blocked_8**. It is finding the **best basis** for the dataset under the sparsity prior:

- DIV2K (transform-needy): L1 picks 6 rotations on the inner-block qubits. PSNR equal to the training-rate-optimal **blocked_b**, structurally simpler (6 gates vs 12, no CP entanglement).
- QuickDraw (already pixel-sparse): L1 picks the identity basis. PSNR far above any **blocked_b**, because no transform was needed.

This is substantially stronger than the original “auto-discover block size” framing: L1 generalises beyond the **blocked_b** family entirely. The block-aligned mask is one valid sparse-basis shape; L1 finds others when they exist (DIV2K's no-CP basin) or returns to identity when no transform is needed (QuickDraw).

The headline **qft < blocked_8** gap is then explained as **Adam-from-qft_identity** is **poorly coupled to the sparse local minima of the QFT(8, 8) loss surface**. Any sparsity-inducing prior — matched mask (block-aware L2), generic sparsity (L1), or even the warm-start init — steers training to one of those minima.

Reproducibility

```
python experiments/qft_identity_regularization.py \
  --gpu 0 --reg block --lambdas 0,1e-3,1e-2,1e-1,1,10 \
  --outer-weight 10 --epochs 112
```

Per-run output: `results/qft_identity_init/div2k_8q/_runs/reg_lambda_<lam>_W<W>/. Run time:  $\approx 9$  min per  $\lambda$  on RTX 3090. 17 unit tests in tests/test_identity_reg.py.`