

Progressive gate-unfreezing of the QFT operator

One gate thawed per stage, trained to a plateau — three unfreeze orderings $\times$ two initialisations on DIV2K-8q

Generated 2026-05-27


Question

The `QFTBasis(8, 8)` (72 $U(2)$ gates) starts either at the QFT-family **identity** ($T_{t=0}(x) = x$) or **Haar-random**. We **unfreeze one gate per stage**, train it to a plateau, then thaw the next. Does the **order** — block-growth (bg), left $\rightarrow$ right (lr), right $\rightarrow$ left (rl) — or the **initialisation** change the trajectory or the endpoint?


QFT gate-unfreezing — three orderings (single axis, $m = 5$)

Same QFT(5) circuit in the paper's tensor-network notation (controlled-phase M = box between two wires with legs to a dot on each). Bold badge = **unfreeze step** for that ordering; fill = block-stage k (= highest qubit a gate touches). Cumulative: each step thaws one more gate, the rest stay frozen at identity. (No SWAP gates — bit-reversal is folded into the index convention.)


(a) **block-growth** stage k completes QFT(k, k); long-range couplings delayed — the gate-level mirror of the block sweep



(b) **left → right sweep** unfreeze in QFT construction order 1→15; q_1 couples to all of q_2 – q_5 before H_2 at step 6



(c) **right → left sweep** reverse construction order; last-built gate (H_5) first, H_1 last



Block-stage colours: ■ $k = 1$ ■ $k = 2$ ■ $k = 3$ ■ $k = 4$ ■ $k = 5$

Figure 1: The three orderings on a QFT circuit (badge = unfreeze step, fill = block-stage k). **bg**: stage k completes QFT(k, k). **lr**: QFT construction order. **rl**: reverse. The same index sequences (at $m = 8$) drive every run.

Setup

Fixed batch of 50 training images; **MSELoss** on the top-10% coefficients. pdft’s JIT Adam step runs with a per-stage frozen set (moments reset per stage); a Riemannian grad-norm probe gives the plateau signal. A stage ends at $\|g\| < 10^{-5}$ **or** $|\Delta L| < 10^{-5}$ (min 5 steps), else an 800-step cap. Random init: Haar $U(2)$ per Hadamard, uniform phase per controlled-phase, seed shared across orderings.

Training dynamics




Figure 2: Absolute top- k MSE vs cumulative step, one curve per ordering (rows: identity / random init). Endpoints carry the final MSE and PSNR@ $\rho\{=\}.20$; tick labels mark **which gate** (e.g. H7, CP3,1) was thawed at each of the largest drops. Per-cell loss+grad-norm staircases: [div2k_8q/<init>/figures/](#).

Reconstruction PSNR (DIV2K)

Test PSNR (dB) at the final stage, with cumulative steps and plateau-trigger counts (grad-norm / loss- Δ / cap):

init	order	$\rho\{=\}.05$	$\rho\{=\}.10$	$\rho\{=\}.15$	$\rho\{=\}.20$	steps (g/ Δ /cap)
identity	block-growth	19.4	26.9	29.7	32	7810 (0/72/0)
identity	left→right	25.2	27.9	30	31.8	7635 (0/72/0)
identity	right→left	24.8	27.2	29.1	30.8	7709 (0/72/0)
random	block-growth	25.2	27.8	29.8	31.7	9962 (0/72/0)
random	left→right	25.2	27.8	29.8	31.7	8415 (0/72/0)
random	right→left	24.8	27.2	29.1	30.8	9904 (0/72/0)

Block-size baseline (comparison)

qft_progressive reaches the **same** QFT(8,8) on DIV2K via a **block-size** curriculum (train all $k(k+1)$ gates of a $2^k \times 2^k$ QFT at once, $k = 1..8$):

k	block	# gates	$\rho\{=\}.05$	$\rho\{=\}.10$	$\rho\{=\}.15$	$\rho\{=\}.20$
1	2×2	2	9.4	12.5	16.2	21.1
2	4×4	6	19.1	26.9	29.8	32
3	8×8	12	25.2	28.1	30.3	32.3
4	16×16	20	25.3	28	30.1	32
5	32×32	30	25.2	27.8	29.8	31.7
6	64×64	42	25.2	27.8	29.8	31.7
7	128×128	56	25.2	27.8	29.8	31.7
8	256×256	72	25.2	27.8	29.8	31.7

It saturates by $k = 2$; the full QFT(8,8) endpoint is **31.7 dB** @ $\rho\{=\}.20$. The gate-unfreeze sweep lands at the same value (identity `bg/lr` 32/31.8 dB, random ≈ 31.7), so the QFT(8,8) optimum is **schedule-independent**. For reference, the best **any** prior trained basis reaches on DIV2K is 33.7 dB (`#pastmax.div2k_8q.by_rho.at("0.20").method`); the QFT family sits ≈ 2 dB under the richer $U(4)$ families.

Reading

The staircase shows each gate's marginal value: a tall drop is a gate that mattered (mostly the early, small-block gates — see the marked steps), a flat plateau a gate left near where it was thawed. All three orderings and both inits converge to ≈ 31.7 dB, so **order and curriculum set only the path, not the destination** — `rl` merely trails by ≈ 1 dB en route. Identity reaches a lower training MSE than random but the same PSNR endpoint.