

Seed robustness of the gate-unfreeze QFT

100 random-initialisation seeds $\times$ three unfreeze orderings, top-20% training objective, DIV2K-8q

Generated 2026-06-22

Question

Does the gate-unfreeze QFT endpoint depend on the random draw? We retrain the `QFTBasis(8,8)` operator from **Haar-random** init under the block-growth (**bg**), left $\rightarrow$ right (**lr**) and right $\rightarrow$ left (**rl**) unfreeze orderings, for 100 seeds each. Every seed s reseeds **everything trainable** — the Haar gate init, the 50-image training-batch subsample (from the fixed 500-image pool), and the training RNG — while the held-out **test set is held fixed** (canonical seed 42) so the endpoint-PSNR spread reflects the trained operator, not the test draw. Training minimises the top-20%-coefficient MSE (the headline elsewhere is top-10%); test PSNR is scored at four keep ratios ρ .

The random initialisations are genuinely different




Figure 1: **(a)** the untrained random operator’s initial top- k MSE loss L_0 on a common fixed batch, per seed — it spans 387.3–2283.1 (mean 1060.8, $\sigma = 417.1$), so the seeds start from genuinely different points. **(b)** a 2-D PCA of the 100 init parameter vectors; the top two components explain only 2.8% + 2.6% of the variance, i.e. the inits are spread across many near-orthogonal directions rather than clustered.

Endpoint variance




Figure 2: Per-ordering test PSNR across 100 seeds. **Top:** (a) mean $\pm\sigma$ band (+ min–max whiskers) vs ρ , against block-FFT 8×8 (dashed) and block-DCT 8×8 (dotted); and (c) the endpoint distribution at $\rho=\}.20$ (histogram + fitted normal). **Bottom:** (b) per-seed scatter at four keep ratios ($\rho = 0.01, 0.05, 0.10, 0.20$) — the trained bases beat block-FFT at every rate and exceed both classical references by ~ 6 dB at $\rho=\}.01$, while block-DCT overtakes them for $\rho \geq .05$.

ordering	$\rho=\}.05$	$\rho=\}.10$	$\rho=\}.15$	$\rho=\}.20$	min@.20
block-growth	25.25±0.03	27.87±0.08	29.94±0.12	31.78±0.16	31.65
left→right	25.22±0.04	27.82±0.05	29.87±0.07	31.7±0.1	31.29
right→left	25.02±0.15	27.53±0.21	29.51±0.26	31.28±0.3	30.16
block-FFT 8×8	24.47	27.1	29.06	30.79	—
block-DCT 8×8	26.11	29.41	31.86	34.01	—

mean $\pm\sigma$ test PSNR (dB), $n = 100$ random-init seeds per ordering; **min@.20** is the single worst seed at $\rho=\}.20$.

Reading

Despite starting from 100 genuinely different random inits (previous section), each ordering’s endpoint sits in a narrow band: at $\rho=\}.20$ the per-seed standard deviation is block-growth 0.16, left→right 0.1, right→left 0.3 dB. The **release order**, not the random draw, is the dominant factor: block-growth and left→right are tight and sit highest, while right→left is $\sim 2\times$ wider and lowest — the dynamics figure below explains why (rl starts from a far worse loss and its seeds stay spread until the final unfreeze stages).

Against the classical block-FFT 8×8 reference (30.79 dB @ $\rho=\}.20$): block-growth and left→right clear it for **every** seed (worst seed 31.65 / 31.29 dB; margins +0.86 / +0.5 dB). right→left clears it for 97 of 100 seeds; 3 low-tail seeds dip below by at most 0.63 dB. Overall 297 of 300 runs (99%) clear block-FFT. So the gate-unfreeze QFT is **essentially always** above block-FFT — unconditionally for bg/lr, and for 97% of the highest-variance

rl ordering — while block-DCT 8×8 (34.01 dB) remains the strongest classical transform here, ahead of the QFT family.

The per-seed endpoints are **not** Gaussian: a Shapiro–Wilk test rejects normality for all three orderings ($p < 10^{-7}$), because seeds settle into a few discrete attractor basins of the very-flat top- k MSE valley rather than scattering smoothly. The fitted normal in the endpoint panel is a visual summary of location and spread, not a claim of Gaussianity.




Figure 3: Per-seed training dynamics: every seed’s top- k MSE staircase (one per unfreeze stage), one panel per ordering, on a shared y-range. Different Haar starts descend along different paths: block-growth (a) collapses fast and stays tight, whereas right→left (c) starts far higher and its seeds remain spread until the final stages — visually accounting for the endpoint variances above. left→right (b) is intermediate.

Emergent block structure of the trained QFT

The same 100-seed operators let us test, robustly across random init, the paper’s central structural claim (sec5.3): the trained full-image QFT factors itself into a block code with no block prior. We read each converged QFTBasis(8,8) directly. Of the 16 Hadamard-role gates (8 per dimension), about 4 per dimension collapse to non-mixing **Pauli-Z/X** gates (mixing $2|a||b| \approx 0.3\%$) — classical block indices — while all 56 controlled-phase gates stay fully active, so the intra-block transform remains rich. The dense operator is consequently block-diagonal: its off-block energy is only 0.002% at a 16-pixel block — essentially exact — with a modal effective block of 16×16 .




Figure 4: Per-Hadamard freeze probability across all seeds, one row per unfreeze ordering (row qubits r1..8 | col qubits c1..8). Roughly half of each dimension’s Hadamards reliably freeze to Pauli-Z/X (non-mixing), the rest stay frequency-mixing.




Figure 5: **Left:** $\log|W|$ of the representative seed’s 1-D operator factor, block-diagonal at the emergent 16-pixel block. **Right:** block-leakage vs candidate block size — the knee at $b = 16$ (dotted) shows the operator is block-structured at that specific scale and not below; the untrained closed-form QFT (not shown) stays high at every scale.

The same factorization is visible directly in frequency space.



Figure 6: Mean test-set power spectrum (peak-normalized $\log_{10} \overline{|F|^2}$, averaged over the 50 DIV2K test images) of the untrained **global QFT** (left) and the **trained QFT** (right, representative seed). The global transform packs energy into a single low-frequency lobe; the trained transform tiles into a 16×16 grid of repeated sub-spectra — the block-periodic signature of a block transform, the same structure the 8×8 block references produce in Fig. 4 of the main text.

Across all 300 operators the structure is stable: the per-seed block-leakage stays negligible (below 0.01%) and the pooled modal effective block is 16×16 , so the block factorization is a robust property of the trained QFT rather than a one-run artifact. The tightly-converged block-growth and left→right orderings concentrate at

16×16 , while the high-variance right $\rightarrow$ left ordering freezes fewer column Hadamards (the lighter c5/c6 cells in the freeze map above), giving larger, more variable column blocks — the one axis along which the emergent block loosens.