

top- k truncation loss $\mathcal{L}_k$
(training batch)

—●— at init (perturbed) exact init
-■- after training

